

Web Spreadsheet

Audrey Tang

This chapter introduces a web spreadsheet written in 99 lines of the three languages natively supported by web browsers: HTML, JavaScript, and CSS.

The ES5 version of this project is available as a jsFiddle¹.

19.1 Introduction

When Tim Berners-Lee invented the web in 1990, *web pages* were written in HTML by marking up text with angle-bracketed *tags*, assigning a logical structure to the content. Text marked up within `<a> . . . ` became *hyperlinks* that would refer the user to other pages on the web.

In the 1990s, browsers added various presentational tags to the HTML vocabulary, including some notoriously nonstandard tags such as `<blink> . . . </blink>` from Netscape Navigator and `<marquee> . . . </marquee>` from Internet Explorer, causing widespread problems in usability and browser compatibility.

In order to restrict HTML to its original purpose—describing a document’s logical structure—browser makers eventually agreed to support two additional languages: CSS to describe presentational styles of a page, and JavaScript (JS) to describe its dynamic interactions.

Since then, the three languages have become more concise and powerful through twenty years of co-evolution. In particular, improvements in JS engines made it practical to deploy large-scale JS frameworks, such as AngularJS².

Today, cross-platform *web applications* (such as web spreadsheets) are as ubiquitous and popular as platform-specific applications (such as VisiCalc, Lotus 1-2-3 and Excel) from the previous century.

How many features can a web application offer in 99 lines with AngularJS? Let’s see it in action!

19.2 Overview

The spreadsheet³ directory contains our showcase for late-2014 editions of the three web languages: HTML⁴ for structure, CSS⁵ for presentation, and the JS ES6 “Harmony”⁶ standard for interaction.

¹<http://jsfiddle.net/audreyt/LtDyP/>

²<http://angularjs.org/>

³<https://github.com/audreyt/500lines/tree/master/spreadsheet/code>

⁴<http://www.w3.org/TR/html5/>

⁵<http://www.w3.org/TR/css3-ui/>

⁶<http://git.io/es6features>

It also uses web storage⁷ for data persistence and web workers⁸ for running JS code in the background. As of this writing, these web standards are supported by Firefox, Chrome, and Internet Explorer 11+, as well as mobile browsers on iOS 5+ and Android 4+.

Now let's open our spreadsheet⁹ in a browser (Figure 19.1):

	A	B	C	D	E
1	1874	+	2046	⇒	3920
2					

Figure 19.1: Initial Screen

Basic Concepts

The spreadsheet spans two dimensions, with *columns* starting from **A**, and *rows* starting from **1**. Each *cell* has a unique *coordinate* (such as **A1**) and *content* (such as “1874”), which belongs to one of four *types*:

- Text: “+” in **B1** and “->” in **D1**, aligned to the left.
- Number: “1874” in **A1** and “2046” in **C1**, aligned to the right.
- Formula: =A1+C1 in **E1**, which *calculates* to the *value* “3920”, displayed with a light blue background.
- Empty: All cells in row **2** are currently empty.

Click “3920” to set *focus* on **E1**, revealing its formula in an *input box* (Figure 19.2).

	A	B	C	D	E
1	1874	+	2046	⇒	=A1+C1
2					

Figure 19.2: Input Box

Now let's set focus on **A1** and *change* its content to “1”, causing **E1** to *recalculate* its value to “2047” (Figure 19.3).

	A	B	C	D	E
1	1	+	2046	⇒	2047
2					

Figure 19.3: Changed Content

Press **ENTER** to set focus to **A2** and change its content to =Date(), then press **TAB**, change the content of **B2** to =alert(), then press **TAB** again to set focus to **C2** (Figure 19.4).

This shows that a formula may calculate to a number (“2047” in **E1**), a text (the current time in **A2**, aligned to the left), or an *error* (red letters in **B2**, aligned to the center).

Next, let's try entering =for(;;){}, the JS code for an infinite loop that never terminates. The spreadsheet will prevent this by automatically *restoring* the content of **C2** after an attempted change.

Now reload the page in the browser with **Ctrl-R** or **Cmd-R** to verify that the spreadsheet content is *persistent*, staying the same across browser sessions. To *reset* the spreadsheet to its original contents, press the ‘curved arrow’ button on the top-left corner.

⁷<http://www.whatwg.org/specs/web-apps/current-work/multipage/webstorage.html>

⁸<http://www.whatwg.org/specs/web-apps/current-work/multipage/workers.html>

⁹<http://audreyt.github.io/500lines/spreadsheet/>

	A	B	C	D	E
1	1	+	2046	⇒	2047
2	Fri Jul 04 2014 21:14:09 GMT+0800 (CST)	ReferenceError: alert is not defined			

Figure 19.4: Formula Error

Progressive Enhancement

Before we dive into the 99 lines of code, it's worthwhile to disable JS in the browser, reload the page, and note the differences (Figure 19.5).

- Instead of a large grid, only a 2x2 table remains onscreen, with a single content cell.
- Row and column labels are replaced by `{{ row }}` and `{{ col }}`.
- Pressing the reset button produces no effect.
- Pressing **TAB** or clicking into the first line of content still reveals an editable input box.

	<code>{{ col }}</code>
<code>{{ row }}</code>	<div> <input type="text"/> </div> <div> <code>errs[col+row]</code> <code> </code> <code>vals[col+row]</code> <code>}}</code> </div>

Figure 19.5: With JavaScript Disabled

When we disable the dynamic interactions (JS), the content structure (HTML) and the presentational styles (CSS) remain in effect. If a website is useful with both JS and CSS disabled, we say it adheres to the *progressive enhancement* principle, making its content accessible to the largest audience possible.

Because our spreadsheet is a web application with no server-side code, we must rely on JS to provide the required logic. However, it does work correctly when CSS is not fully supported, such as with screen readers and text-mode browsers.

	A	B	C	D	E
1	1874	+	2046	⇒	=A1+C1
	1874	+	2046	⇒	3920
2					

Figure 19.6: With CSS Disabled

As shown in Figure 19.6, if we enable JS in the browser and disable CSS instead, the effects are:

- All background and foreground colors are gone.
- The input box and the cell value are both displayed, instead of just one at a time.
- Otherwise, the application still works the same as the full version.

19.3 Code Walkthrough

Figure 19.7 shows the links between HTML and JS components. In order to make sense of the diagram, let's go through the four source code files, in the same sequence as the browser loads them.

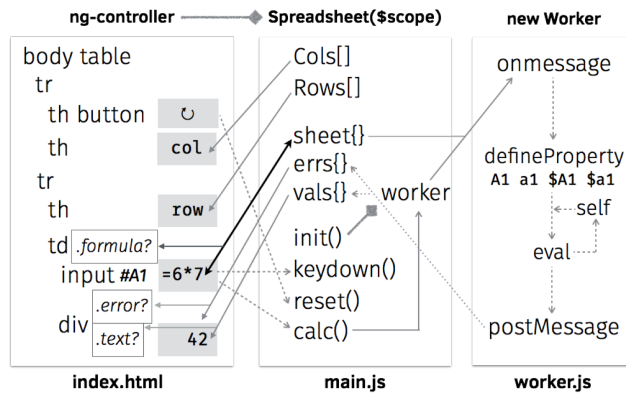


Figure 19.7: Architecture Diagram

- **index.html**: 19 lines
- **main.js**: 38 lines (excluding comments and blank lines)
- **worker.js**: 30 lines (excluding comments and blank lines)
- **styles.css**: 12 lines

HTML

The first line in `index.html` declares that it's written in HTML5 with the UTF-8 encoding:

```
<!DOCTYPE html><html><head><meta charset="UTF-8">
```

Without the charset declaration, the browser may display the reset button's Unicode symbol as `â†`, an example of *mojibake*: garbled text caused by decoding issues.

The next three lines are JS declarations, placed within the head section as usual:

```
<script src="lib/angular.js"></script>
<script src="main.js"></script>
<script>
  try { angular.module('500lines') }
  catch(e){ location="es5/index.html" }
</script>
```

The `<script src=". . .">` tags load JS resources from the same path as the HTML page. For example, if the current URL is `http://abc.com/x/index.html`, then `lib/angular.js` refers to `http://abc.com/x/lib/angular.js`.

The `try{ angular.module('500lines') }` line tests if `main.js` is loaded correctly; if not, it tells the browser to navigate to `es5/index.html` instead. This *redirect-based graceful degradation* technique ensures that for pre-2015 browsers with no ES6 support, we can use the translated-to-ES5 versions of JS programs as a fallback.

The next two lines load the CSS resource, close the head section, and begin the body section containing the user-visible part:

```
<link href="styles.css" rel="stylesheet">
</head><body ng-app="500lines" ng-controller="Spreadsheet" ng-cloak>
```

The `ng-app` and `ng-controller` attributes above tell AngularJS¹⁰ to call the `500lines` module's `Spreadsheet` function, which would return a *model*: an object that provides *bindings* on the document *view*. (The `ng-cloak` attribute hides the document from display until the bindings are in place.)

As a concrete example, when the user clicks the `<button>` defined in the next line, its `ng-click` attribute will trigger and call `reset()` and `calc()`, two named functions provided by the JS model:

```
<table><tr>
  <th><button type="button" ng-click="reset(); calc()"></button></th>
```

The next line uses `ng-repeat` to display the list of column labels on the top row:

```
<th ng-repeat="col in Cols">{{ col }}</th>
```

For example, if the JS model defines `Cols` as `["A", "B", "C"]`, then there will be three heading cells (`th`) labeled accordingly. The `{{ col }}` notation tells AngularJS to *interpolate* the expression, filling the contents in each `th` with the current value of `col`.

Similarly, the next two lines go through values in `Rows` — `[1, 2, 3]` and so on — creating a row for each one and labeling the leftmost `th` cell with its number:

```
</tr><tr ng-repeat="row in Rows">
  <th>{{ row }}</th>
```

Because the `<tr ng-repeat>` tag is not yet closed by `</tr>`, the `row` variable is still available for expressions. The next line creates a data cell (`td`) in the current row and uses both `col` and `row` variables in its `ng-class` attribute:

```
<td ng-repeat="col in Cols" ng-class="{ formula: ('=' === sheet[col+row][0]) }">
```

A few things are going on here. In HTML, the `class` attribute describes a *set of class names* that allow CSS to style them differently. The `ng-class` here evaluates the expression `('=' === sheet[col+row][0])`; if it is true, then the `<td>` gets `formula` as an additional class, which gives the cell a light-blue background as defined in line 8 of **styles.css** with the `.formula class selector`.

The expression above checks if the current cell is a formula by testing if `=` is the initial character (`[0]`) of the string in `sheet[col+row]`, where `sheet` is a JS model object with coordinates (such as `"E1"`) as properties, and cell contents (such as `"=A1+C1"`) as values. Note that because `col` is a string and not a number, the `+` in `col+row` means concatenation instead of addition.

Inside the `<td>`, we give the user an input box to edit the cell content stored in `sheet[col+row]`:

```
<input id="{{ col+row }}" ng-model="sheet[col+row]" ng-change="calc()"
  ng-model-options="{ debounce: 200 }" ng-keydown="keydown( $event, col, row )">
```

¹⁰<http://angularjs.org/>

Here, the key attribute is `ng-model`, which enables a *two-way binding* between the JS model and the input box's editable content. In practice, this means that whenever the user makes a change in the input box, the JS model will update `sheet[col+row]` to match the content, and trigger its `calc()` function to recalculate values of all formula cells.

To avoid repeated calls to `calc()` when the user presses and holds a key, `ng-model-options` limits the update rate to once every 200 milliseconds.

The `id` attribute here is interpolated with the coordinate `col+row`. The `id` attribute of a HTML element must be different from the `id` of all other elements in the same document. This ensures that the `#A1` ID selector refers to a single element, instead of a set of elements like the class selector `.formula`. When the user presses the **UP/DOWN/ENTER** keys, the keyboard-navigation logic in `keydown()` will use ID selectors to determine which input box to focus on.

After the input box, we place a `<div>` to display the calculated value of the current cell, represented in the JS model by objects `errs` and `vals`:

```
<div ng-class="{ error: errs[col+row], text: vals[col+row][0] }">
  {{ errs[col+row] || vals[col+row] }}</div>
```

If an error occurs when computing a formula, the text interpolation uses the error message contained in `errs[col+row]`, and `ng-class` applies the error class to the element, allowing CSS to style it differently (with red letters, aligned to the center, etc.).

When there is no error, the `vals[col+row]` on the right side of `||` is interpolated instead. If it's a non-empty string, the initial character (`[0]`) will evaluate to true, applying the `text` class to the element that left-aligns the text.

Because empty strings and numeric values have no initial character, `ng-class` will not assign them any classes, so CSS can style them with right alignment as the default case.

Finally, we close the `ng-repeat` loop in the column level with `</td>`, close the row-level loop with `</tr>`, and end the HTML document with:

```
</td>
</tr></table>
</body></html>
```

JS: Main Controller

The `main.js` file defines the `500lines` module and its Spreadsheet controller function, as required by the `<body>` element in `index.html`.

As the bridge between the HTML view and the background worker, it has four tasks:

- Define the dimensions and labels of columns and rows.
- Provide event handlers for keyboard navigation and the reset button.
- When the user changes the spreadsheet, send its new content to the worker.
- When computed results arrive from the worker, update the view and save the current state.

The flowchart in Figure 19.8 shows the controller-worker interaction in more detail:

Now let's walk through the code. In the first line, we request the AngularJS `$scope`:

```
angular.module('500lines', []).controller('Spreadsheet', function ($scope, $timeout) {
```

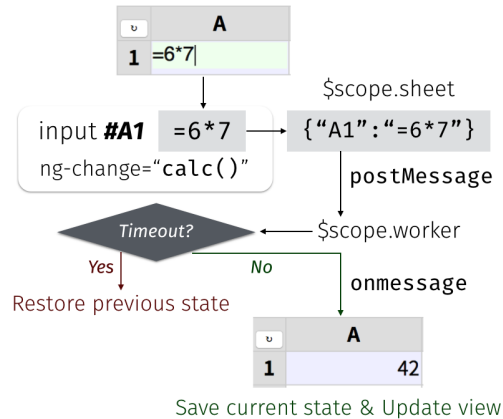


Figure 19.8: Controller-Worker Flowchart

The \$ in \$scope is part of the variable name. Here we also request the \$timeout¹¹ service function from AngularJS; later on, we will use it to prevent infinite-looping formulas.

To put Cols and Rows into the model, simply define them as properties of \$scope:

```
// Begin of $scope properties; start with the column/row labels
$scope.Cols = [], $scope.Rows = [];
for (col of range( 'A', 'H' )) { $scope.Cols.push(col); }
for (row of range( 1, 20 )) { $scope.Rows.push(row); }
```

The ES6 for...of¹² syntax makes it easy to loop through ranges with a start and an end point, with the helper function range defined as a generator¹³:

```
function* range(cur, end) { while (cur <= end) { yield cur;
```

The function* above means that range returns an iterator¹⁴, with a while loop that would yield¹⁵ a single value at a time. Whenever the for loop demands the next value, it will resume execution right after the yield line:

```
// If it's a number, increase it by one; otherwise move to next letter
cur = (isNaN( cur ) ? String.fromCharCode( cur.codePointAt()+1 ) : cur+1);
} }
```

To generate the next value, we use isNaN to see if cur is meant as a letter (NaN stands for “not a number.”) If so, we get the letter’s code point value¹⁶, increment it by one, and convert the codepoint¹⁷ back to get its next letter. Otherwise, we simply increase the number by one.

¹¹[https://docs.angularjs.org/api/ng/service/\\$timeout](https://docs.angularjs.org/api/ng/service/$timeout)

¹²<https://developer.mozilla.org/en-US/docs/Web/JavaScript/Reference/Statements/for...of>

¹³https://developer.mozilla.org/en-US/docs/Web/JavaScript/Reference/Statements/function*

¹⁴https://developer.mozilla.org/en-US/docs/Web/JavaScript/Guide/The_iterator_protocol

¹⁵<https://developer.mozilla.org/en-US/docs/Web/JavaScript/Reference/Operators/yield>

¹⁶https://developer.mozilla.org/en-US/docs/Web/JavaScript/Reference/Global_Objects/String/codePointAt

¹⁷https://developer.mozilla.org/en-US/docs/Web/JavaScript/Reference/Global_Objects/String/fromCodePoint

Next up, we define the `keydown()` function that handles keyboard navigation across rows:

```
// UP(38) and DOWN(40)/ENTER(13) move focus to the row above (-1) and below (+1).
$scope.keydown = ({which}, col, row)=>{ switch (which) {
```

The arrow function¹⁸ receives the arguments (`$event`, `col`, `row`) from `<input ng-keydown>`, using destructuring assignment¹⁹ to assign `$event.which` into the `which` parameter, and checks if it's among the three navigational key codes:

```
case 38: case 40: case 13: $timeout( )=>{
```

If it is, we use `$timeout` to schedule a focus change after the current `ng-keydown` and `ng-change` handler. Because `$timeout` requires a function as argument, the `()=>{ . . . }` syntax constructs a function to represent the focus-change logic, which starts by checking the direction of movement:

```
const direction = (which === 38) ? -1 : +1;
```

The `const` declarator means `direction` will not change during the function's execution. The direction to move is either upward (-1, from **A2** to **A1**) if the key code is 38 (**UP**), or downward (+1, from **A2** to **A3**) otherwise.

Next up, we retrieve the target element using the ID selector syntax (e.g. `"#A3"`), constructed with a template string²⁰ written in a pair of backticks, concatenating the leading #, the current `col` and the target `row + direction`:

```
const cell = document.querySelector( `#${ col }${ row + direction }` );
if (cell) { cell.focus(); }
} );
} };
```

We put an extra check on the result of `querySelector` because moving upward from **A1** will produce the selector `#A0`, which has no corresponding element, and so will not trigger a focus change — the same goes for pressing **DOWN** at the bottom row.

Next, we define the `reset()` function so the reset button can restore the contents of the sheet:

```
// Default sheet content, with some data cells and one formula cell.
$scope.reset = ()=>{
  $scope.sheet = { A1: 1874, B1: '+', C1: 2046, D1: '->', E1: '=A1+C1' }; }
```

The `init()` function tries restoring the sheet content from its previous state from the `localStorage`²¹, and defaults to the initial content if it's our first time running the application:

```
// Define the initializer, and immediately call it
($scope.init = ()=>{
  // Restore the previous .sheet; reset to default if it's the first run
  $scope.sheet = angular.fromJson( localStorage.getItem( '' ) );
  if (!$scope.sheet) { $scope.reset(); }
  $scope.worker = new Worker( 'worker.js' );
}).call();
```

¹⁸https://developer.mozilla.org/en-US/docs/Web/JavaScript/Reference/arrow_functions

¹⁹https://developer.mozilla.org/en-US/docs/Web/JavaScript/New_in_JavaScript/1.7\#Pulling_fields_from_objects_passed_as_function_parameter

²⁰https://developer.mozilla.org/en-US/docs/Web/JavaScript/Reference/template_strings

²¹<https://developer.mozilla.org/en-US/docs/Web/Guide/API/DOM/Storage\#localStorage>

A few things are worth noting in the `init()` function above:

- We use the `($scope.init = ()=>{ . . . }).call()` syntax to define the function and immediately call it.
- Because `localStorage` only stores strings, we *parse* the sheet structure from its JSON²² representation using `angular.fromJson()`.
- At the last step of `init()`, we create a new web worker²³ thread and assign it to the `worker` scope property. Although the worker is not directly used in the view, it's customary to use `$scope` to share objects used across model functions, in this case between `init()` here and `calc()` below.

While `sheet` holds the user-editable cell content, `errs` and `vals` contain the results of calculations — errors and values — that are read-only to the user:

```
// Formula cells may produce errors in .errs; normal cell contents are in .vals
[$scope.errs, $scope.vals] = [ {}, {} ];
```

With these properties in place, we can define the `calc()` function that triggers whenever the user makes a change to `sheet`:

```
// Define the calculation handler; not calling it yet
$scope.calc = ()=>{
  const json = angular.toJson( $scope.sheet );
```

Here we take a snapshot of the state of `sheet` and store it in the constant `json`, a JSON string. Next up, we construct a promise from `$timeout`²⁴ that cancels the upcoming computation if it takes more than 99 milliseconds:

```
const promise = $timeout( ()=>{
  // If the worker has not returned in 99 milliseconds, terminate it
  $scope.worker.terminate();
  // Back up to the previous state and make a new worker
  $scope.init();
  // Redo the calculation using the last-known state
  $scope.calc();
}, 99 );
```

Since we made sure that `calc()` is called at most once every 200 milliseconds via the `<input ng-model-options>` attribute in HTML, this arrangement leaves 101 milliseconds for `init()` to restore `sheet` to the last known-good state and make a new worker.

The worker's task is to calculate `errs` and `vals` from the contents of `sheet`. Because **main.js** and **worker.js** communicate by message-passing, we need an `onmessage` handler to receive the results once they are ready:

```
// When the worker returns, apply its effect on the scope
$scope.worker.onmessage = ({data})=>{
  $timeout.cancel( promise );
  localStorage.setItem( '', json );
  $timeout( ()=>{ [$scope.errs, $scope.vals] = data; } );
};
```

²²<https://developer.mozilla.org/en-US/docs/Glossary/JSON>

²³<https://developer.mozilla.org/en-US/docs/Web/API/Worker>

²⁴[https://docs.angularjs.org/api/ng/service/\\$timeout](https://docs.angularjs.org/api/ng/service/$timeout)

If `onmessage` is called, we know that the sheet snapshot in `json` is stable (i.e., containing no infinite-looping formulas), so we cancel the 99-millisecond timeout, write the snapshot to local-Storage, and schedule a UI update with a `$timeout` function that updates `errs` and `vals` to the user-visible view.

With the handler in place, we can post the state of sheet to the worker, starting its calculation in the background:

```
// Post the current sheet content for the worker to process
$scope.worker.postMessage( $scope.sheet );
};

// Start calculation when worker is ready
$scope.worker.onmessage = $scope.calc;
$scope.worker.postMessage( null );
});
```

JS: Background Worker

There are three reasons for using a web worker to calculate formulas, instead of using the main JS thread for the task:

- While the worker runs in the background, the user is free to continue interacting with the spreadsheet without getting blocked by computation in the main thread.
- Because we accept any JS expression in a formula, the worker provides a *sandbox* that prevents formulas from interfering with the page that contains them, such as by popping out an `alert()` dialog box.
- A formula can refer to any coordinates as variables. The other coordinates may contain another formula that might end in a cyclic reference. To solve this problem, we use the worker's *global scope* object `self`, and define these variables as *getter functions* on `self` to implement the cycle-prevention logic.

With these in mind, let's take a look at the worker's code.

The worker's sole purpose is defining its `onmessage` handler. The handler takes `sheet`, calculates `errs` and `vals`, and posts them back to the main JS thread. We begin by re-initializing the three variables when we receive a message:

```
let sheet, errs, vals;
self.onmessage = ({data})=>{
  [sheet, errs, vals] = [ data, {}, {} ];
```

In order to turn coordinates into global variables, we first iterate over each property in `sheet`, using a `for...in` loop:

```
for (const coord in sheet) {
```

ES6 introduces `const` and `let` declares *block scoped* constants and variables; `const coord` above means that functions defined in the loop would capture the value of `coord` in each iteration.

In contrast, `var coord` in earlier versions of JS would declare a *function scoped* variable, and functions defined in each loop iteration would end up pointing to the same `coord` variable.

Customarily, formula variables are case-insensitive and can optionally have a `$` prefix. Because JS variables are case-sensitive, we use `map` to go over the four variable names for the same coordinate:

```
// Four variable names pointing to the same coordinate: A1, a1, $A1, $a1
[ '', '$' ].map( p => [ coord, coord.toLowerCase() ].map(c => {
  const name = p+c;
```

Note the shorthand arrow function syntax above: `p => ...` is the same as `(p) => { ... }`.

For each variable name, like `A1` and `$a1`, we define an accessor property²⁵ on `self` that calculates `vals["A1"]` whenever they are evaluated in an expression:

```
// Worker is reused across calculations, so only define each variable once
if ((Object.getOwnPropertyDescriptor( self, name ) || {}).get) { return; }

// Define self['A1'], which is the same thing as the global variable A1
Object.defineProperty( self, name, { get() {
```

The `{ get() { ... } }` syntax above is shorthand for `{ get: ()=>{ ... } }`. Because we define only `get` and not `set`, the variables become *read-only* and cannot be modified from user-supplied formulas.

The `get` accessor starts by checking `vals[coord]`, and simply returns it if it's already calculated:

```
if (coord in vals) { return vals[coord]; }
```

If not, we need to calculate `vals[coord]` from `sheet[coord]`.

First we set it to `NaN`, so self-references like setting **A1** to =A1 will end up with `NaN` instead of an infinite loop:

```
vals[coord] = NaN;
```

Next we check if `sheet[coord]` is a number by converting it to numeric with prefix `+`, assigning the number to `x`, and comparing its string representation with the original string. If they differ, then we set `x` to the original string:

```
// Turn numeric strings into numbers, so =A1+C1 works when both are numbers
let x = +sheet[coord];
if (sheet[coord] !== x.toString()) { x = sheet[coord]; }
```

If the initial character of `x` is `=`, then it's a formula cell. We evaluate the part after `=` with `eval.call()`, using the first argument `null` to tell `eval` to run in the *global scope*, hiding the *lexical scope* variables like `x` and `sheet` from the evaluation:

```
// Evaluate formula cells that begin with =
try { vals[coord] = (('=' === x[0]) ? eval.call( null, x.slice( 1 ) ) : x);
```

If the evaluation succeeds, the result is stored into `vals[coord]`. For non-formula cells, the value of `vals[coord]` is simply `x`, which may be a number or a string.

If `eval` results in an error, the catch block tests if it's because the formula refers to an empty cell not yet defined in `self`:

```
} catch (e) {
  const match = /\$?[A-Za-z][1-9][0-9]*\b/.exec( e );
  if (match && !( match[0] in self )) {
```

²⁵https://developer.mozilla.org/en-US/docs/Web/JavaScript/Reference/Global_Objects/Object/defineProperty

In that case, we set the missing cell's default value to "0", clear `vals[coord]`, and re-run the current computation using `self[coord]`:

```
    // The formula refers to a uninitialized cell; set it to 0 and retry
    self[match[0]] = 0;
    delete vals[coord];
    return self[coord];
}
```

If the user gives the missing cell a content later on in `sheet[coord]`, then the temporary value would be overridden by `Object.defineProperty`.

Other kinds of errors are stored in `errs[coord]`:

```
    // Otherwise, stringify the caught exception in the errs object
    errs[coord] = e.toString();
}
```

In case of errors, the value of `vals[coord]` will remain NaN because the assignment did not finish executing.

Finally, the get accessor returns the calculated value stored in `vals[coord]`, which must be a number, a Boolean value, or a string:

```
    // Turn vals[coord] into a string if it's not a number or Boolean
    switch (typeof vals[coord]) {
      case 'function': case 'object': vals[coord]+='';
    }
    return vals[coord];
  } } );
});
}
```

With accessors defined for all coordinates, the worker goes through the coordinates again, invoking each accessor with `self[coord]`, then posts the resulting `errs` and `vals` back to the main JS thread:

```
// For each coordinate in the sheet, call the property getter defined above
for (const coord in sheet) { self[coord]; }
return [ errs, vals ];
}
```

CSS

The **styles.css** file contains just a few selectors and their presentational styles. First, we style the table to merge all cell borders together, leaving no spaces between neighboring cells:

```
table { border-collapse: collapse; }
```

Both the heading and data cells share the same border style, but we can tell them apart by their background colors: heading cells are light gray, data cells are white by default, and formula cells get a light blue background:

```
th, td { border: 1px solid #ccc; }
th { background: #ddd; }
td.formula { background: #eef; }
```

The displayed width is fixed for each cell's calculated values. Empty cells receive a minimal height, and long lines are clipped with a trailing ellipsis:

```
td div { text-align: right; width: 120px; min-height: 1.2em;
        overflow: hidden; text-overflow: ellipsis; }
```

The text alignment and decorations are determined by each value's type, as reflected by the text and error class selectors:

```
div.text { text-align: left; }
div.error { text-align: center; color: #800; font-size: 90%; border: solid 1px #800 }
```

As for the user-editable input box, we use *absolute positioning* to overlay it on top of its cell, and make it transparent so the underlying div with the cell's value shows through:

```
input { position: absolute; border: 0; padding: 0;
        width: 120px; height: 1.3em; font-size: 100%;
        color: transparent; background: transparent; }
```

When the user sets focus on the input box, it springs into the foreground:

```
input:focus { color: #111; background: #efe; }
```

Furthermore, the underlying div is collapsed into a single line, so it's completely covered by the input box:

```
input:focus + div { white-space: nowrap; }
```

19.4 Conclusion

Since this book is *500 Lines or Less*, a web spreadsheet in 99 lines is a minimal example—please feel free to experiment and extend it in any direction you'd like.

Here are some ideas, all easily reachable in the remaining space of 401 lines:

- A collaborative online editor using ShareJS²⁶, AngularFire²⁷ or GoAngular²⁸.
- Markdown syntax support for text cells, using angular-marked²⁹.
- Common formula functions (SUM, TRIM, etc.) from the OpenFormula standard³⁰.
- Interoperate with popular spreadsheet formats, such as CSV and SpreadsheetML via SheetJS³¹.
- Import from and export to online spreadsheet services, such as Google Spreadsheet and EtherCalc³².

²⁶<http://sharejs.org/>

²⁷<http://angularfire.com>

²⁸<http://goangular.org/>

²⁹<http://ngmodules.org/modules/angular-marked>

³⁰<https://en.wikipedia.org/wiki/OpenFormula>

³¹<http://sheetjs.com/>

³²<http://ethercalc.net/>

A Note on JS versions

This chapter aims to demonstrate new concepts in ES6, so we use the Traceur compiler³³ to translate source code to ES5 to run on pre-2015 browsers.

If you prefer to work directly with the 2010 edition of JS, the `as-javascript-1.8.5`³⁴ directory has **main.js** and **worker.js** written in the style of ES5; the source code³⁵ is line-by-line comparable to the ES6 version with the same line count.

For people preferring a cleaner syntax, the `as-livescript-1.3.0`³⁶ directory uses LiveScript³⁷ instead of ES6 to write **main.ls** and **worker.ls**; it is 20 lines shorter³⁸ than the JS version.

Building on the LiveScript language, the `as-react-livescript`³⁹ directory uses the ReactJS⁴⁰ framework; it is 10 lines more longer⁴¹ than the AngularJS equivalent, but runs considerably faster.

If you are interested in translating this example to alternate JS languages, send a pull request⁴²—I'd love to hear about it!

³³<https://github.com/google/traceur-compiler>

³⁴<https://audreyt.github.io/500lines/spreadsheet/as-javascript-1.8.5/>

³⁵<https://github.com/audreyt/500lines/tree/master/spreadsheet/as-javascript-1.8.5>

³⁶<https://audreyt.github.io/500lines/spreadsheet/as-livescript-1.3.0/>

³⁷<http://livescript.net/>

³⁸<https://github.com/audreyt/500lines/tree/master/spreadsheet/as-livescript-1.3.0>

³⁹<https://audreyt.github.io/500lines/spreadsheet/as-react-livescript/>

⁴⁰<https://facebook.github.io/react/>

⁴¹<https://github.com/audreyt/500lines/tree/master/spreadsheet/as-react-livescript>

⁴²<https://github.com/audreyt/500lines/pulls>

Static Analysis

Leah Hanson

20.1 Introduction

You may be familiar with a fancy IDE that draws red underlines under parts of your code that don't compile. You may have run a linter on your code to check for formatting or style problems. You might run your compiler in super-picky mode with all the warnings turned on. All of these tools are applications of static analysis.

Static analysis is a way to check for problems in your code without running it. “Static” means at compile time rather than at run time, and “analysis” means we’re analyzing the code. When you’ve used the tools I mentioned above, it may have felt like magic. But those tools are just programs—they are made of source code that was written by a person, a programmer like you. In this chapter, we’re going to talk about how to implement a couple of static analysis checks. In order to do this, we need to know what we want the check to do and how we want to do it.

We can get more specific about what you need to know by describing the process as three stages:

1. Deciding what you want to check for.

You should be able to explain the general problem you’d like to solve, in terms that a user of the programming language would recognize. Examples include:

- Finding misspelled variable names
- Finding race conditions in parallel code
- Finding calls to unimplemented functions

2. Deciding how exactly to check for it.

While we could ask a friend to do one of the tasks listed above, they aren’t specific enough to explain to a computer. To tackle “misspelled variable names”, for example, we’d need to decide what misspelled means here. One option would be to claim variable names should be composed of English words from the dictionary; another option is to look for variables that are only used once (the one time you mistyped it).

If we know we’re looking for variables that are only used once, we can talk about kinds of variable usages (having their value assigned versus read) and what code would or would not trigger a warning.

3. Implementation details.

This covers the actual act of writing the code, the time spent reading the documentation for libraries you use, and figuring out how to get at the information you need to write the analysis. This could involve reading in a file of code, parsing it to understand the structure, and then making your specific check on that structure.

We’re going to work through these steps for each of the individual checks implemented in this chapter. Step 1 requires enough understanding of the language we’re analyzing to empathize with the problems its users face. All the code in this chapter is Julia code, written to analyze Julia code.

20.2 A Very Brief Introduction to Julia

Julia is a young language aimed at technical computing. It was released at version 0.1 in the spring of 2012; as of the start of 2015, it has reached version 0.3. In general, Julia looks a lot like Python, but with some optional type annotations and without any object-oriented stuff. The feature that most programmers will find novel in Julia is multiple dispatch, which has a pervasive impact on both API design and on other design choices in the language.

Here is a snippet of Julia code:

```
# A comment about increment
function increment(x::Int64)
    return x + 1
end

increment(5)
```

This code defines a method of the function `increment` that takes one argument, named `x`, of type `Int64`. The method returns the value of `x + 1`. Then, this freshly defined method is called with the value 5; the function call, as you may have guessed, will evaluate to 6.

`Int64` is a type whose values are signed integers represented in memory by 64 bits; they are the integers that your hardware understands if your computer has a 64-bit processor. Types in Julia define the representation of data in memory, in addition to influencing method dispatch.

The name `increment` refers to a generic function, which may have many methods. We have just defined one method of it. In many languages, the terms “function” and “method” are used interchangeably; in Julia, they have distinct meanings. This chapter will make more sense if you are careful to understand “function” as a named collection of methods, where a “method” is a specific implementation for a specific type signature.

Let’s define another method of the `increment` function:

```
# Increment x by y
function increment(x::Int64, y::Number)
    return x + y
end

increment(5) # => 6
increment(5,4) # => 9
```

Now the function `increment` has two methods. Julia decides which method to run for a given call based on the number and types of the arguments; this is called *dynamic multiple dispatch*:

- **Dynamic** because it's based on the types of the values used at runtime.
- **Multiple** because it looks at the types and order of all the arguments.
- **Dispatch** because this is a way of matching function calls to method definitions.

To put this in the context of languages you may already know, object-oriented languages use single dispatch because they only consider the first argument. (In `x.foo(y)`, the first argument is `x`.)

Both single and multiple dispatch are based on the types of the arguments. The `x::Int64` above is a type annotation purely for dispatch. In Julia's dynamic type system, you could assign a value of any type to `x` during the function without an error.

We haven't really seen the "multiple" part yet, but if you're curious about Julia, you'll have to look that up on your own. We need to move on to our first check.

20.3 Checking the Types of Variables in Loops

As in most programming languages, writing very fast code in Julia involves an understanding of how the computer works and how Julia works. An important part of helping the compiler create fast code for you is writing type-stable code; this is important in Julia and JavaScript, and is also helpful in other JIT'd languages. When the compiler can see that a variable in a section of code will always contain the same specific type, the compiler can do more optimizations than if it believes (correctly or not) that there are multiple possible types for that variable. You can read more about why type stability (also called "monomorphism") is important for JavaScript online¹.

Why This Is Important

Let's write a function that takes an `Int64` and increases it by some amount. If the number is small (less than 10), let's increase it by a big number (50), but if it's big, let's only increase it by 0.5.

```
function increment(x::Int64)
    if x < 10
        x = x + 50
    else
        x = x + 0.5
    end
    return x
end
```

This function looks pretty straightforward, but the type of `x` is unstable. I selected two numbers: 50, an `Int64`, and 0.5, a `Float64`. Depending on the value of `x`, it might be added to either one of them. If you add an `Int64` like 22, to a `Float64` like 0.5, you'll get a `Float64` (22.5). Because the type of variable in the function (`x`) could change depending on the value of the arguments to the function (`x`), this method of `increment` and specifically the variable `x` are type-unstable.

`Float64` is a type that represents floating-point values stored in 64 bits; in C, it is called a double. This is one of the floating-point types that 64-bit processors understand.

As with most efficiency problems, this issue is more pronounced when it happens during loops. Code inside for loops and while loops is run many, many times, so making it fast is more important than speeding up code that is only run once or twice. Therefore, our first check is to look for variables that have unstable types inside loops.

¹<http://mrale.ph/blog/2015/01/11/whats-up-with-monomorphism.html>

First, let's look at an example of what we want to catch. We'll be looking at two functions. Each of them sums the numbers 1 to 100, but instead of summing the whole numbers, they divide each one by 2 before summing it. Both functions will get the same answer (2525.0); both will return the same type (Float64). However, the first function, `unstable`, suffers from type-instability, while the second one, `stable`, does not.

```
function unstable()
    sum = 0
    for i=1:100
        sum += i/2
    end
    return sum
end

function stable()
    sum = 0.0
    for i=1:100
        sum += i/2
    end
    return sum
end
```

The only textual difference between the two functions is in the initialization of `sum`: `sum = 0` versus `sum = 0.0`. In Julia, `0` is an `Int64` literal and `0.0` is a `Float64` literal. How big of a difference could this tiny change make?

Because Julia is Just-In-Time (JIT) compiled, the first run of a function will take longer than subsequent runs. (The first run includes the time it takes to compile the function for these argument types.) When we benchmark functions, we have to be sure to run them once (or precompile them) before timing them.

```
julia> unstable()
2525.0

julia> stable()
2525.0

julia> @time unstable()
elapsed time: 9.517e-6 seconds (3248 bytes allocated)
2525.0

julia> @time stable()
elapsed time: 2.285e-6 seconds (64 bytes allocated)
2525.0
```

The `@time` macro prints out how long the function took to run and how many bytes were allocated while it was running. The number of bytes allocated increases every time new memory is needed; it does not decrease when the garbage collector vacuums up memory that's no longer being used. This means that the bytes allocated is related to the amount of time we spend allocating and managing memory, but does not imply that we had all of that memory in use at the same time.

If we wanted to get solid numbers for `stable` versus `unstable` we would need to make the loop much longer or run the functions many times. However, it looks like `unstable` is probably slower.

More interestingly, we can see a large gap in the number of bytes allocated; `unstable` has allocated around 3 KB of memory, where `stable` is using 64 bytes.

Since we can see how simple `unstable` is, we might guess that this allocation is happening in the loop. To test this, we can make the loop longer and see if the allocations increase accordingly. Let's make the loop go from 1 to 10000, which is 100 times more iterations; we'll look for the number of bytes allocated to also increase about 100 times, to around 300 KB.

```
function unstable()
    sum = 0
    for i=1:10000
        sum += i/2
    end
    return sum
end
```

Since we redefined the function, we'll need to run it so it gets compiled before we measure it. We expect to get a different, larger answer from the new function definition, since it's summing more numbers now.

```
julia> unstable()
2.50025e7

julia>@time unstable()
elapsed time: 0.000667613 seconds (320048 bytes allocated)
2.50025e7
```

The new `unstable` allocated about 320 KB, which is what we would expect if the allocations are happening in the loop. To explain what's going on here, we're going to look at how Julia works under the hood.

This difference between `unstable` and `stable` occurs because `sum` in `unstable` must be boxed while `sum` in `stable` can be unboxed. Boxed values consist of a type tag and the actual bits that represent the value; unboxed values only have their actual bits. But the type tag is small, so that's not why boxing values allocates a lot more memory.

The difference comes from what optimizations the compiler can make. When a variable has a concrete, immutable type, the compiler can unbox it inside the function. If that's not the case, then the variable must be allocated on the heap, and participate in the garbage collector. Immutable types are a concept specific to Julia. A value of an immutable type can't be changed.

Immutable types are usually types that represent values, rather than collections of values. For example, most numeric types, including `Int64` and `Float64`, are immutable. (Numeric types in Julia are normal types, not special primitive types; you could define a new `MyInt64` that's the same as the provided one.) Because immutable types cannot be modified, you must make a new copy every time you want change one. For example `4 + 6` must make a new `Int64` to hold the result. In contrast, the members of a mutable type can be updated in-place; this means you don't have to make a copy of the whole thing to make a change.

The idea of `x = x + 2` allocating memory probably sounds pretty weird; why would you make such a basic operation slow by making `Int64` values immutable? This is where those compiler optimizations come in: using immutable types doesn't (usually) slow this down. If `x` has a stable, concrete type (such as `Int64`), then the compiler is free to allocate `x` on the stack and mutate `x` in place. The problem is only when `x` has an unstable type (so the compiler doesn't know how big or

what type it will be); once `x` is boxed and on the heap, the compiler isn't completely sure that some other piece of code isn't using the value, and thus can't edit it.

Because `sum` in `stable` has a concrete type (`Float64`), the compiler knows that it can store it unboxed locally in the function and mutate its value; `sum` will not be allocated on the heap and new copies don't have to be made every time we add `i/2`.

Because `sum` in `unstable` does not have a concrete type, the compiler allocates it on the heap. Every time we modify `sum`, we allocated a new value on the heap. All this time spent allocating values on the heap (and retrieving them every time we want to read the value of `sum`) is expensive.

Using `0` versus `0.0` is an easy mistake to make, especially when you're new to Julia. Automatically checking that variables used in loops are type-stable helps programmers get more insight into what the types of their variables are in performance-critical sections of their code.

Implementation Details

We'll need to find out which variables are used inside loops and we'll need to find the types of those variables. We'll then need to decide how to print them in a human-readable format.

- How do we find loops?
- How do we find variables in loops?
- How do we find the types of a variable?
- How do we print the results?
- How do we tell if the type is unstable?

I'm going to tackle the last question first, since this whole endeavour hinges on it. We've looked at an unstable function and seen, as programmers, how to identify an unstable variable, but we need our program to find them. This sounds like it would require simulating the function to look for variables whose values might change—which sounds like it would take some work. Luckily for us, Julia's type inference already traces through the function's execution to determine the types.

The type of `sum` in `unstable` is `Union{Float64, Int64}`. This is a `UnionType`, a special kind of type that indicates that the variable may hold any of a set of types of values. A variable of type `Union{Float64, Int64}` can hold values of type `Int64` or `Float64`; a value can only have one of those types. A `UnionType` joins any number of types (e.g., `UnionType{Float64, Int64, Int32}` joins three types). We're going to look for is `UnionTyped` variables inside loops.

Parsing code into a representative structure is a complicated business, and gets more complicated as the language grows. In this chapter, we'll be depending on internal data structures used by the compiler. This means that we don't have to worry about reading files or parsing them, but it does mean we have to work with data structures that are not in our control and that sometimes feel clumsy or ugly.

Besides all the work we'll save by not having to parse the code by ourselves, working with the same data structures that the compiler uses means that our checks will be based on an accurate assessment of the compiler's understanding—which means our check will be consistent with how the code actually runs.

This process of examining Julia code from Julia code is called introspection. When you or I introspect, we're thinking about how and why we think and feel. When code introspects, it examines the representation or execution properties of code in the same language (possibly its own code). When code's introspection extends to modifying the examined code, it's called metaprogramming (programs that write or modify programs).

Introspection in Julia

Julia makes it easy to introspect. There are four functions built in to let us see what the compiler is thinking: `code_lowered`, `code_typed`, `code_llvm`, and `code_native`. Those are listed in order of what step in the compilation process their output is from; the first one is closest to the code we'd type in and the last one is the closest to what the CPU runs. For this chapter, we'll focus on `code_typed`, which gives us the optimized, type-inferred abstract syntax tree (AST).

`code_typed` takes two arguments: the function of interest, and a tuple of argument types. For example, if we wanted to see the AST for a function `foo` when called with two `Int64`s, then we would call `code_typed(foo, (Int64,Int64))`.

```
function foo(x,y)
    z = x + y
    return 2 * z
end

code_typed(foo, (Int64,Int64))
```

This is the structure that `code_typed` would return:

```
1-element Array{Any,1}:
:($Expr(:lambda, { :x,:y}, { { :z}, { { :x,Int64,0}, { :y,Int64,0}, { :z,Int64,18} }, {} },
:(begin # none, line 2:
    z = (top(box))(Int64,(top(add_int))(x::Int64,y::Int64))::Int64 # line 3:
    return (top(box))(Int64,(top(mul_int))(2,z::Int64))::Int64
end::Int64)))
```

This is an `Array`; this allows `code_typed` to return multiple matching methods. Some combinations of functions and argument types may not completely determine which method should be called. For example, you could pass in a type like `Any` (instead of `Int64`). `Any` is the type at the top of the type hierarchy; all types are subtypes of `Any` (including `Any`). If we included `Any` in our tuple of argument types, and had multiple matching methods, then the `Array` from `code_typed` would have more than one element in it; it would have one element per matching method.

Let's pull our example `Expr` out to make it easier to talk about.

```
julia> e = code_typed(foo, (Int64,Int64))[1]
:($Expr(:lambda, { :x,:y}, { { :z}, { { :x,Int64,0}, { :y,Int64,0}, { :z,Int64,18} }, {} },
:(begin # none, line 2:
    z = (top(box))(Int64,(top(add_int))(x::Int64,y::Int64))::Int64 # line 3:
    return (top(box))(Int64,(top(mul_int))(2,z::Int64))::Int64
end::Int64)))
```

The structure we're interested in is inside the `Array`: it is an `Expr`. Julia uses `Expr` (short for expression) to represent its AST. (An abstract syntax tree is how the compiler thinks about the meaning of your code; it's kind of like when you had to diagram sentences in grade school.) The `Expr` we get back represents one method. It has some metadata (about the variables that appear in the method) and the expressions that make up the body of the method.

Now we can ask some questions about `e`.

We can ask what properties an `Expr` has by using the `names` function, which works on any Julia value or type. It returns an `Array` of names defined by that type (or the type of the value).

```
julia> names(e)
3-element Array{Symbol,1}:
 :head
 :args
 :typ
```

We just asked `e` what names it has, and now we can ask what value each name corresponds to. An `Expr` has three properties: `head`, `typ` and `args`.

```
julia> e.head
:lambda
```

```
julia> e.typ
Any
```

```
julia> e.args
3-element Array{Any,1}:
{:x,:y}
{{:z},{{:x,Int64,0},{:y,Int64,0},{:z,Int64,18}},{}}
:(begin # none, line 2:
    z = (top(box))(Int64,(top(add_int))(x::Int64,y::Int64))::Int64 # line 3:
    return (top(box))(Int64,(top(mul_int))(2,z::Int64))::Int64
end)::Int64)
```

We just saw some values printed out, but that doesn't tell us much about what they mean or how they're used.

- `head` tells us what kind of expression this is; normally, you'd use separate types for this in Julia, but `Expr` is a type that models the structure used in the parser. The parser is written in a dialect of Scheme, which structures everything as nested lists. `head` tells us how the rest of the `Expr` is organized and what kind of expression it represents.
- `typ` is the inferred return type of the expression; when you evaluate any expression, it results in some value. `typ` is the type of the value that the expression will evaluate to. For nearly all `Expr`s, this value will be `Any` (which is always correct, since every possible type is a subtype of `Any`). Only the body of type-inferred methods and most expressions inside them will have their `typ` set to something more specific. (Because `type` is a keyword, this field can't use that word as its name.)
- `args` is the most complicated part of `Expr`; its structure varies based on the value of `head`. It's always an `Array{Any}` (an untyped array), but beyond that the structure changes.

In an `Expr` representing a method, there will be three elements in `e.args`:

```
julia> e.args[1] # names of arguments as symbols
2-element Array{Any,1}:
 :x
 :y
```

Symbols are a special type for representing the names of variables, constants, functions, and modules. They are a different type from strings because they specifically represent the name of a program construct.

```
julia> e.args[2] # three lists of variable metadata
3-element Array{Any,1}:
{:z}
{:x,Int64,0},{:y,Int64,0},{:z,Int64,18}}
{}
```

The first list above contains the names of all local variables; we only have one (z) here. The second list contains a tuple for each variable in and argument to the method; each tuple has the variable name, the variable's inferred type, and a number. The number conveys information about how the variable is used, in a machine- (rather than human-) friendly way. The last list is of captured variable names; it's empty in this example.

```
julia> e.args[3] # the body of the method
:(begin # none, line 2:
    z = (top(box))(Int64,(top(add_int))(x::Int64,y::Int64))::Int64 # line 3:
    return (top(box))(Int64,(top(mul_int))(2,z::Int64))::Int64
end::Int64)
```

The first two args elements are metadata about the third. While the metadata is very interesting, it isn't necessary right now. The important part is the body of the method, which is the third element. This is another Expr.

```
julia> body = e.args[3]
:(begin # none, line 2:
    z = (top(box))(Int64,(top(add_int))(x::Int64,y::Int64))::Int64 # line 3:
    return (top(box))(Int64,(top(mul_int))(2,z::Int64))::Int64
end::Int64)
```

```
julia> body.head
:body
```

This Expr has head :body because it's the body of the method.

```
julia> body.typ
Int64
```

The typ is the inferred return type of the method.

```
julia> body.args
4-element Array{Any,1}:
:( # none, line 2:)
:(z = (top(box))(Int64,(top(add_int))(x::Int64,y::Int64))::Int64)
:( # line 3:)
:(return (top(box))(Int64,(top(mul_int))(2,z::Int64))::Int64)
```

args holds a list of expressions: the list of expressions in the method's body. There are a couple of annotations of line numbers (i.e., :(# line 3:)), but most of the body is setting the value of z (z = x + y) and returning 2 * z. Notice that these operations have been replaced by Int64-specific intrinsic functions. The top(function-name) indicates an intrinsic function; something that is implemented in Julia's code generation, rather than in Julia.

We haven't seen what a loop looks like yet, so let's try that.

```

julia> function lloop(x)
    for x = 1:100
        x *= 2
    end
end
lloop (generic function with 1 method)

julia> code_typed(lloop, (Int,))[1].args[3]
:(begin # none, line 2:
    #s120 = $(Expr(:new, UnitRange{Int64}, 1, :(((top(getfield))(Intrinsics,
    :select_value))((top(sle_int))(1,100)::Bool,100,(top(box))(Int64,(top(
    sub_int))(1,1))::Int64)::Int64)))::UnitRange{Int64}
    #s119 = (top(getfield))(#s120::UnitRange{Int64},:start)::Int64          unless
    (top(box))(Bool,(top(not_int))(#s119::Int64 == (top(box))(Int64,(top(
    add_int))((top(getfield))
    (#s120::UnitRange{Int64},:stop)::Int64,1))::Int64::Bool))::Bool goto 1
2:
    _var0 = #s119::Int64
    _var1 = (top(box))(Int64,(top(add_int))(#s119::Int64,1))::Int64
    x = _var0::Int64
    #s119 = _var1::Int64 # line 3:
    x = (top(box))(Int64,(top(mul_int))(x::Int64,2))::Int64
3:
    unless (top(box))(Bool,(top(not_int))((top(box))(Bool,(top(not_int))
    (#s119::Int64 == (top(box))(Int64,(top(add_int))((top(getfield))
    #s120::UnitRange{Int64},:stop)::Int64,1))::Int64::Bool))::Bool)
        goto 2
1:
        0:
    return
end::Nothing)

```

You'll notice there's no for or while loop in the body. As the compiler transforms the code from what we wrote to the binary instructions the CPU understands, features that are useful to humans but that are not understood by the CPU (like loops) are removed. The loop has been rewritten as label and goto expressions. The goto has a number in it; each label also has a number. The goto jumps to the the label with the same number.

Detecting and Extracting Loops

We're going to find loops by looking for goto expressions that jump backwards.

We'll need to find the labels and gotos, and figure out which ones match. I'm going to give you the full implementation first. After the wall of code, we'll take it apart and examine the pieces.

```

# This is a function for trying to detect loops in the body of a Method
# Returns lines that are inside one or more loops
function loopcontents(e::Expr)
    b = body(e)
    loops = Int[]
    nesting = 0
    lines = {}
    for i in 1:length(b)

```



```

if typeof(b[i]) == LabelNode
    l = b[i].label
    jumpback = findnext(x-> (typeof(x) == GotoNode && x.label == l)
                        || (Base.is_expr(x,:gotoifnot) && x.args[end] == l),
                        b, i)
    if jumpback != 0
        push!(loops,jumpback)
        nesting += 1
    end
end
if nesting > 0
    push!(lines,(i,b[i]))
end

if typeof(b[i]) == GotoNode && in(i,loops)
    splice!(loops,findfirst(loops,i))
    nesting -= 1
end
end
lines
end

```

And now to explain in pieces:

`b = body(e)`

We start by getting all the expressions in the body of method, as an Array. `body` is a function that I've already implemented:

```

# Return the body of a Method.
# Takes an Expr representing a Method,
# returns Vector{Expr}.
function body(e::Expr)
    return e.args[3].args
end

```

And then:

```

loops = Int[]
nesting = 0
lines = {}

```

`loops` is an Array of label line numbers where gotos that are loops occur. `nesting` indicates the number of loops we are currently inside. `lines` is an Array of (index, Expr) tuples.

```

for i in 1:length(b)
    if typeof(b[i]) == LabelNode
        l = b[i].label
        jumpback = findnext(
            x-> (typeof(x) == GotoNode && x.label == l)
              || (Base.is_expr(x,:gotoifnot) && x.args[end] == l),
            b, i)
    end
end

```

```

    if jumpback != 0
        push!(loops,jumpback)
        nesting += 1
    end
end

```

We look at each expression in the body of `e`. If it is a label, we check to see if there is a `goto` that jumps to this label (and occurs after the current index). If the result of `findnext` is greater than zero, then such a `goto` node exists, so we'll add that to `loops` (the Array of loops we are currently in) and increment our nesting level.

```

    if nesting > 0
        push!(lines,(i,b[i]))
    end

```

If we're currently inside a loop, we push the current line to our array of lines to return.

```

    if typeof(b[i]) == GotoNode && in(i,loops)
        splice!(loops,findfirst(loops,i))
        nesting -= 1
    end
end
lines
end

```

If we're at a `GotoNode`, then we check to see if it's the end of a loop. If so, we remove the entry from `loops` and reduce our nesting level.

The result of this function is the `lines` array, an array of (index, value) tuples. This means that each value in the array has an index into the method-body-Expr's body and the value at that index. Each element of `lines` is an expression that occurred inside a loop.

Finding and Typing Variables

We just finished the function `loopcontents` which returns the Exprs that are inside loops. Our next function will be `loosetypes`, which takes a list of Exprs and returns a list of variables that are loosely typed. Later, we'll pass the output of `loopcontents` into `loosetypes`.

In each expression that occurred inside a loop, `loosetypes` searches for occurrences of symbols and their associated types. Variable usages show up as `SymbolNodes` in the AST; `SymbolNodes` hold the name and inferred type of the variable.

We can't just check each expression that `loopcontents` collected to see if it's a `SymbolNode`. The problem is that each Expr may contain one or more Expr; each Expr may contain one or more `SymbolNodes`. This means we need to pull out any nested Exprs, so that we can look in each of them for `SymbolNodes`.

```

# given `lr`, a Vector of expressions (Expr + literals, etc)
# try to find all occurrences of a variables in `lr`
# and determine their types
function loosetypes(lr::Vector)
    symbols = SymbolNode[]
    for (i,e) in lr

```

```

    if typeof(e) == Expr
      es = copy(e.args)
      while !isempty(es)
        e1 = pop!(es)
        if typeof(e1) == Expr
          append!(es,e1.args)
        elseif typeof(e1) == SymbolNode
          push!(symbols,e1)
        end
      end
    end
  end
end
loose_types = SymbolNode[]
for symnode in symbols
  if !isleafnode(symnode.typ) && typeof(symnode.typ) == UnionType
    push!(loose_types, symnode)
  end
end
return loose_types
end

symbols = SymbolNode[]
for (i,e) in lr
  if typeof(e) == Expr
    es = copy(e.args)
    while !isempty(es)
      e1 = pop!(es)
      if typeof(e1) == Expr
        append!(es,e1.args)
      elseif typeof(e1) == SymbolNode
        push!(symbols,e1)
      end
    end
  end
end
end
end

```

The while loop goes through the guts of all the Exprs, recursively. Every time the loop finds a SymbolNode, it adds it to the vector symbols.

```

loose_types = SymbolNode[]
for symnode in symbols
  if !isleafnode(symnode.typ) && typeof(symnode.typ) == UnionType
    push!(loose_types, symnode)
  end
end
return loose_types
end

```

Now we have a list of variables and their types, so it's easy to check if a type is loose. loose_types does that by looking for a specific kind of non-concrete type, a UnionType. We get a lot more “failing” results when we consider all non-concrete types to be “failing”. This is because we’re evaluating each method with its annotated argument types, which are likely to be abstract.

Making This Usable

Now that we can do the check on an expression, we should make it easier to call on a user's code. We'll create two ways to call `checklooptypes`:

1. On a whole function; this will check each method of the given function.
2. On an expression; this will work if the user extracts the results of `code_typed` themselves.

```
## for a given Function, run checklooptypes on each Method
function checklooptypes(f::Callable;kwargs...)
    lrs = LoopResult[]
    for e in code_typed(f)
        lr = checklooptypes(e)
        if length(lr.lines) > 0 push!(lrs,lr) end
    end
    LoopResults(f.env.name,lrs)
end

# for an Expr representing a Method,
# check that the type of each variable used in a loop
# has a concrete type
checklooptypes(e::Expr;kwargs...) =
    LoopResult(MethodSignature(e),loosetypes(loopcontents(e)))
```

We can see both options work about the same for a function with one method:

```
julia> using TypeCheck

julia> function foo(x::Int)
    s = 0
    for i = 1:x
        s += i/2
    end
    return s
end
foo (generic function with 1 method)

julia> checklooptypes(foo)
foo(Int64)::Union{Int64,Float64}
s::Union{Int64,Float64}
s::Union{Int64,Float64}

julia> checklooptypes(code_typed(foo,(Int,))[1])
(Int64)::Union{Int64,Float64}
s::Union{Int64,Float64}
s::Union{Int64,Float64}
```

Pretty Printing

I've skipped an implementation detail here: how did we get the results to print out to the REPL?

First, I made some new types. `LoopResults` is the result of checking a whole function; it has the function name and the results for each method. `LoopResult` is the result of checking one method; it has the argument types and the loosely typed variables.

The `checklooptypes` function returns a `LoopResults`. This type has a function called `show` defined for it. The REPL calls `display` on values it wants to display; `display` will then call our `show` implementation.

This code is important for making this static analysis usable, but it is not doing static analysis. You should use the preferred method for pretty-printing types and output in your implementation language; this is just how it's done in Julia.

```
type LoopResult
    msig::MethodSignature
    lines::Vector{SymbolNode}
    LoopResult(ms::MethodSignature, ls::Vector{SymbolNode}) = new(ms, unique(ls))
end

function Base.show(io::IO, x::LoopResult)
    display(x.msig)
    for snode in x.lines
        println(io, "\t", string(snode.name), ":", string(snode.typ))
    end
end

type LoopResults
    name::Symbol
    methods::Vector{LoopResult}
end

function Base.show(io::IO, x::LoopResults)
    for lr in x.methods
        print(io, string(x.name))
        display(lr)
    end
end
```

20.4 Looking For Unused Variables

Sometimes, as you're typing in your program, you mistype a variable name. The program can't tell that you meant for this to be the same variable that you spelled correctly before; it sees a variable used only one time, where you might see a variable name misspelled. Languages that require variable declarations naturally catch these misspellings, but many dynamic languages don't require declarations and thus need an extra layer of analysis to catch them.

We can find misspelled variable names (and other unused variables) by looking for variables that are only used once—or only used one way.

Here is an example of a little bit of code with one misspelled name.

```
function foo(variable_name::Int)
    sum = 0
    for i=1:variable_name
```

```

    sum += variable_name
end
variable_nme = sum
return variable_name
end

```

This kind of mistake can cause problems in your code that are only discovered when it's run. Let's assume you misspell each variable name only once. We can separate variable usages into writes and reads. If the misspelling is a write (i.e., `worng = 5`), then no error will be thrown; you'll just be silently putting the value in the wrong variable—and it could be frustrating to find the bug. If the misspelling is a read (i.e., `right = worng + 2`), then you'll get a runtime error when the code is run; we'd like to have a static warning for this, so that you can find this error sooner, but you will still have to wait until you run the code to see the problem.

As code becomes longer and more complicated, it becomes harder to spot the mistake—unless you have the help of static analysis.

Left-Hand Side and Right-Hand Side

Another way to talk about “read” and “write” usages is to call them “right-hand side” (RHS) and “left-hand side” (LHS) usages. This refers to where the variable is relative to the `=` sign.

Here are some usages of `x`:

- Left-hand side:

```

- x = 2
- x = y + 22
- x = x + y + 2
- x += 2 (which de-sugars to x = x + 2)

```

- Right-hand side:

```

- y = x + 22
- x = x + y + 2
- x += 2 (which de-sugars to x = x + 2)
- 2 * x
- x

```

Notice that expressions like `x = x + y + 2` and `x += 2` appear in both sections, since `x` appears on both sides of the `=` sign.

Looking for Single-Use Variables

There are two cases we need to look for:

1. Variables used once.
2. Variables used only on the LHS or only on the RHS.

We'll look for all variable usages, but we'll look for LHS and RHS usages separately, to cover both cases.

Finding LHS Usages

To be on the LHS, a variable needs to have an = sign to be to the left of. This means we can look for = signs in the AST, and then look to the left of them to find the relevant variable.

In the AST, an = is an Expr with the head : (=). (The parentheses are there to make it clear that this is the symbol for = and not another operator, :=.) The first value in args will be the variable name on its LHS. Because we're looking at an AST that the compiler has already cleaned up, there will (nearly) always be just a single symbol to the left of our = sign.

Let's see what that means in code:

```
julia> :(x = 5)
:(x = 5)

julia> :(x = 5).head
:(=)

julia> :(x = 5).args
2-element Array{Any,1}:
 :x
 5

julia> :(x = 5).args[1]
:x
```

Below is the full implementation, followed by an explanation.

```
# Return a list of all variables used on the left-hand-side of assignment (=)
#
# Arguments:
#   e: an Expr representing a Method, as from code_typed
#
# Returns:
#   a Set{Symbol}, where each element appears on the LHS of an assignment in e.
#
function find_lhs_variables(e::Expr)
    output = Set{Symbol}()
    for ex in body(e)
        if Base.is_expr(ex, : (=))
            push!(output, ex.args[1])
        end
    end
    return output
end

output = Set{Symbol}()
```

We have a set of Symbols; those are variables names we've found on the LHS.

```
for ex in body(e)
    if Base.is_expr(ex, : (=))
        push!(output, ex.args[1])
    end
end
```

We aren't digging deeper into the expressions, because the `code_typed` AST is pretty flat; loops and ifs have been converted to flat statements with `gotos` for control flow. There won't be any assignments hiding inside function calls' arguments. This code will fail if anything more than a symbol is on the left of the equal sign. This misses two specific edge cases: array accesses (like `a[5]`, which will be represented as a `:ref` expression) and properties (like `a.head`, which will be represented as a `:.expression`). These will still always have the relevant symbol as the first value in their args, it might just be buried a bit (as in `a.property.name.head.other_property`). This code doesn't handle those cases, but a couple lines of code inside the `if` statement could fix that.

```
push!(output,ex.args[1])
```

When we find a LHS variable usage, we push! the variable name into the Set. The Set will make sure that we only have one copy of each name.

Finding RHS usages

To find all the other variable usages, we also need to look at each `Expr`. This is a bit more involved, because we care about basically all the `Exprs`, not just the `:(=)` ones and because we have to dig into nested `Exprs` (to handle nested function calls).

Here is the full implementation, with explanation following.

```
# Given an Expression, finds variables used in it (on right-hand-side)
#
# Arguments: e: an Expr
#
# Returns: a Set{Symbol}, where each e is used in a rhs expression in e
#
function find_rhs_variables(e::Expr)
    output = Set{Symbol}()

    if e.head == :lambda
        for ex in body(e)
            union!(output,find_rhs_variables(ex))
        end
    elseif e.head == :(:=)
        for ex in e.args[2:end] # skip lhs
            union!(output,find_rhs_variables(ex))
        end
    elseif e.head == :return
        output = find_rhs_variables(e.args[1])
    elseif e.head == :call
        start = 2 # skip function name
        e.args[1] == TopNode(:box) && (start = 3) # skip type name
        for ex in e.args[start:end]
            union!(output,find_rhs_variables(ex))
        end
    elseif e.head == :if
        for ex in e.args # want to check condition, too
            union!(output,find_rhs_variables(ex))
        end
    elseif e.head == :(::)
```



```

    output = find_rhs_variables(e.args[1])
end

return output
end

```

The main structure of this function is a large if-else statement, where each case handles a different head-symbol.

```
output = Set{Symbol}()
```

output is the set of variable names, which we will return at the end of the function. Since we only care about the fact that each of these variables has been read at least once, using a Set frees us from worrying about the uniqueness of each name.

```

if e.head == :lambda
    for ex in body(e)
        union!(output, find_rhs_variables(ex))
    end
end

```

This is the first condition in the if-else statement. A :lambda represents the body of a function. We recurse on the body of the definition, which should get all the RHS variable usages in the definition.

```

elseif e.head == :(:=)
    for ex in e.args[2:end] # skip lhs
        union!(output, find_rhs_variables(ex))
    end
end

```

If the head is :(:=), then the expression is an assignment. We skip the first element of args because that's the variable being assigned to. For each of the remaining expressions, we recursively find the RHS variables and add them to our set.

```

elseif e.head == :return
    output = find_rhs_variables(e.args[1])
end

```

If this is a return statement, then the first element of args is the expression whose value is returned; we'll add any variables in there into our set.

```

elseif e.head == :call
    # skip function name
    for ex in e.args[2:end]
        union!(output, find_rhs_variables(ex))
    end
end

```

For function calls, we want to get all variables used in all the arguments to the call. We skip the function name, which is the first element of args.

```

elseif e.head == :if
    for ex in e.args # want to check condition, too
        union!(output, find_rhs_variables(ex))
    end
end

```

An Expr representing an if statement has the head value `:if`. We want to get variable usages from all the expressions in the body of the if statement, so we recurse on each element of `args`.

```
elseif e.head == :(:)
    output = find_rhs_variables(e.args[1])
end
```

The `:(:)` operator is used to add type annotations. The first argument is the expression or variable being annotated; we check for variable usages in the annotated expression.

```
return output
```

At the end of the function, we return the set of RHS variable usages.

There's a little more code that simplifies the method above. Because the version above only handles Exprs, but some of the values that get passed recursively may not be Exprs, we need a few more methods to handle the other possible types appropriately.

```
# Recursive Base Cases, to simplify control flow in the Expr version
find_rhs_variables(a) = Set{Symbol}() # unhandled, should be immediate val e.g. Int
find_rhs_variables(s::Symbol) = Set{Symbol}([s])
find_rhs_variables(s::SymbolNode) = Set{Symbol}([s.name])
```

Putting It Together

Now that we have the two functions defined above, we can use them together to find variables that are either only read from or only written to. The function that finds them will be called `unused_locals`.

```
function unused_locals(e::Expr)
    lhs = find_lhs_variables(e)
    rhs = find_rhs_variables(e)
    setdiff(lhs, rhs)
end
```

`unused_locals` will return a set of variable names. It's easy to write a function that determines whether the output of `unused_locals` counts as a “pass” or not. If the set is empty, the method passes. If all the methods of a function pass, then the function passes. The function `check_locals` below implements this logic.

```
check_locals(f::Callable) = all([check_locals(e) for e in code_typed(f)])
check_locals(e::Expr) = isempty(unused_locals(e))
```

20.5 Conclusion

We've done two static analyses of Julia code—one based on types and one based on variable usages.

Statically-typed languages already do the kind of work our type-based analysis did; additional type-based static analysis is mostly useful in dynamically typed languages. There have been (mostly research) projects to build static type inference systems for languages including Python, Ruby, and Lisp. These systems are usually built around optional type annotations; you can have static types when you want them, and fall back to dynamic typing when you don't. This is especially helpful for integrating some static typing into existing code bases.

Non-typed-based checks, like our variable-usage one, are applicable to both dynamically and statically typed languages. However, many statically typed languages, like C++ and Java, require you to declare variables, and already give basic warnings like the ones we created. There are still custom checks that can be written; for example, checks that are specific to your project's style guide or extra safety precautions based on security policies.

While Julia does have great tools for enabling static analysis, it's not alone. Lisp, of course, is famous for having the code be a data structure of nested lists, so it tends to be easy to get at the AST. Java also exposes its AST, although the AST is much more complicated than Lisp's. Some languages or language tool-chains are not designed to allow mere users to poke around at internal representations. For open-source tool chains (especially well-commented ones), one option is to add hooks to the environment that let you access the AST.

In cases where that won't work, the final fallback is writing a parser yourself; this is to be avoided when possible. It's a lot of work to cover the full grammar of most programming languages, and you'll have to update it yourself as new features are added to the language (rather than getting the updates automatically from upstream). Depending on the checks you want to do, you may be able to get away with parsing only some lines or a subset of language features, which would greatly decrease the cost of writing your own parser.

Hopefully, your new understanding of how static analysis tools are written will help you understand the tools you use on your code, and maybe inspire you to write one of your own.

